

Convolutional Neural Networks

Areeb Gani, Michael Ilie, Vijay Shanmugam


Welcome!



ml.mbhs.edu

Motivation

- Often, with neural networks, we have to deal with images
- Images have lots of data, as they are represented as multiple 2-dimensional arrays




Motivation

- We want to “simplify” the image, keeping its core features, while making it easier for our model to process all the data and perform the normal computations
- We can achieve this using some special functions which we call **kernels**

Convolutional Kernel

- “Sliding window,” called **filter**, applies transformation on certain parts of the image sequentially
- Returns a smaller array than original image, but still retains high-level features (the *gist* of the image)



Convolutional Kernel

- Mathematically, our filter is just a matrix - we multiply the element of the image array with the elements of our matrix to get our output

1 <small>x1</small>	1 <small>x0</small>	1 <small>x1</small>	0	0
0 <small>x0</small>	1 <small>x1</small>	1 <small>x0</small>	1	0
0 <small>x1</small>	0 <small>x0</small>	1 <small>x1</small>	1	1
0	0	1	1	0
0	1	1	0	0

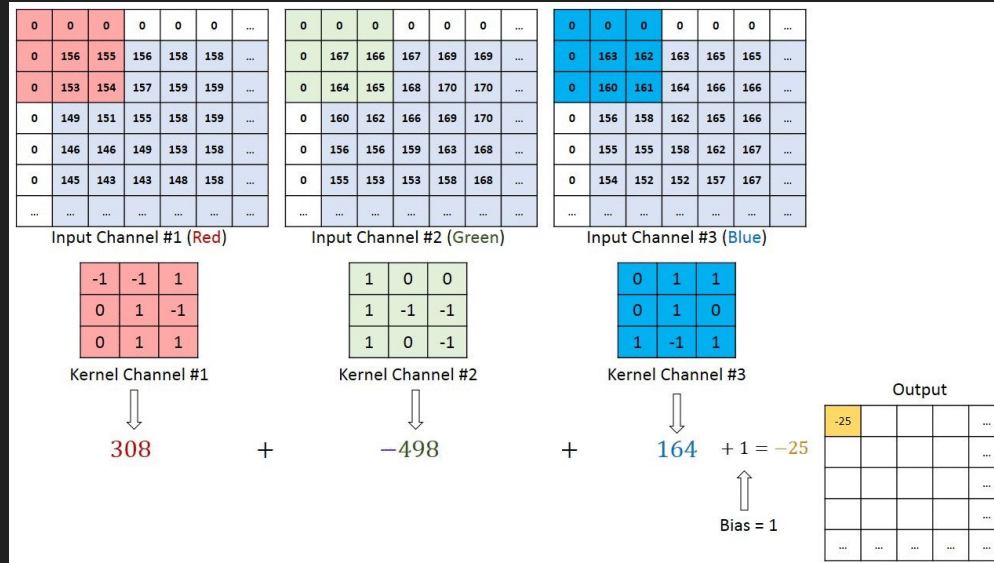
Image

4		

Convolved
Feature

Convolutional Kernel

- For higher dimensions, we add together the output of each filter in each color channel to our final output matrix




Pooling Layer

- After applying the convolution filter, we run our output matrix through a **pooling layer**
- There are different types of pooling layers, but the most common is called “maximum pooling,” which is what we’ll focus on

Pooling Layer

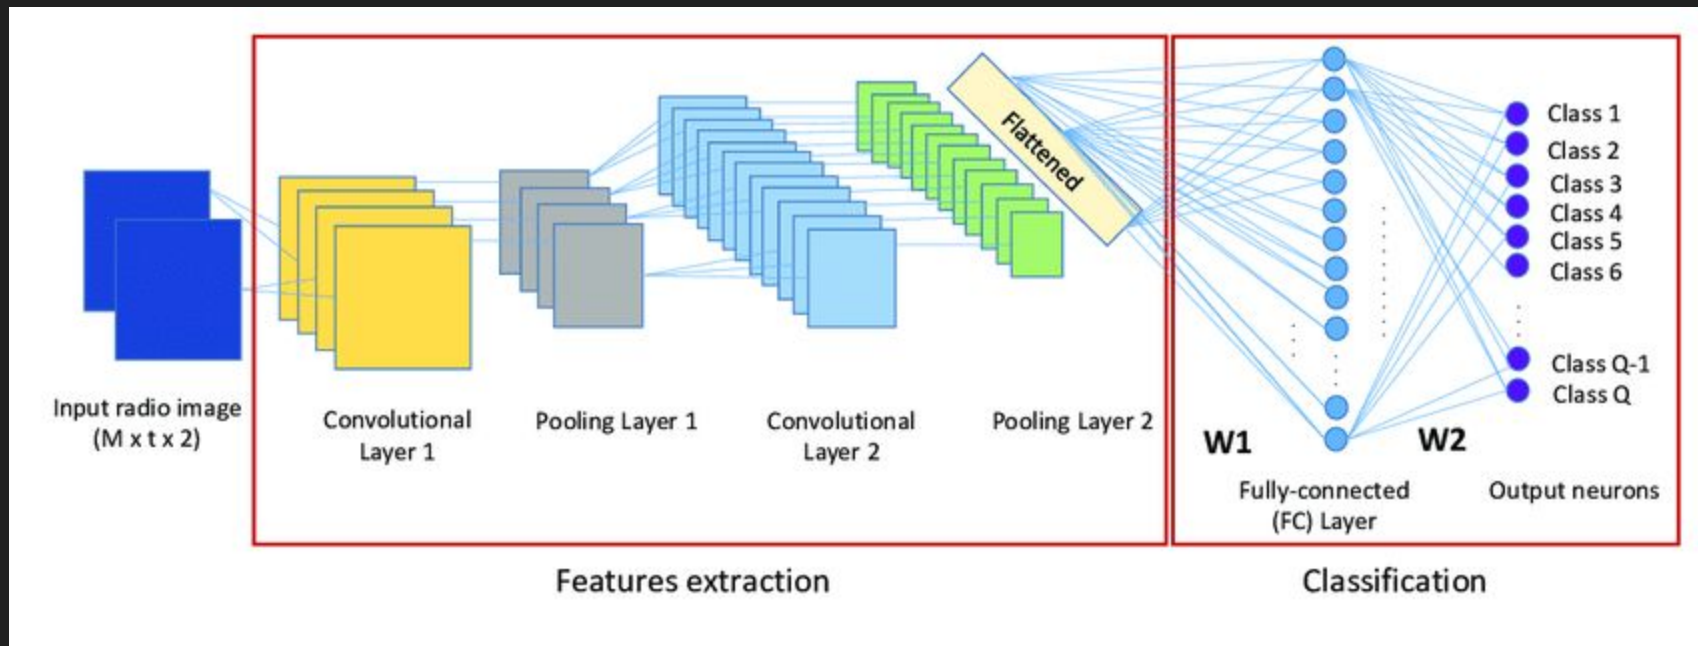
- Max pooling has a similar “sliding window” process to convolutional kernels, but instead of multiplying by a matrix, we simply take the maximum value of each window in our output
- This will find the most dominant features in the image, which is what we need to make our prediction



Fully Connected Layer and Output

- We apply the previous two steps (Convolutional Kernel and Pooling Layer) repeatedly, to minimize the size of the image, for as long as we wish
- Then we use a **fully connected layer** to get our output - this is just a standard neural network layer like we have discussed before
- Finally, we get our output using an activation function such as *sigmoid* or *softmax*

Full Model



Some Practice Tutorials

- <https://www.tensorflow.org/tutorials/images/cnn> (TensorFlow)
- https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html (PyTorch)
- <https://www.analyticsvidhya.com/blog/2019/10/building-image-classification-models-cnn-pytorch/> (also PyTorch)

Join Our Groups

- Sign up for Discord (<https://discord.gg/3Z5YuPqt>)
- Join Deepnote (<https://deepnote.com/join-team?token=af3af0284bc8497>)
- Fill out our form (<https://forms.gle/Fr31aFLWx8cHdtTY8>)
 - Join mailing list + Github organization